

CS 5530



Database Systems Spring 2020

Welcome!


Course Overview

Intro to Databases


Domains of Data




Domains of Data



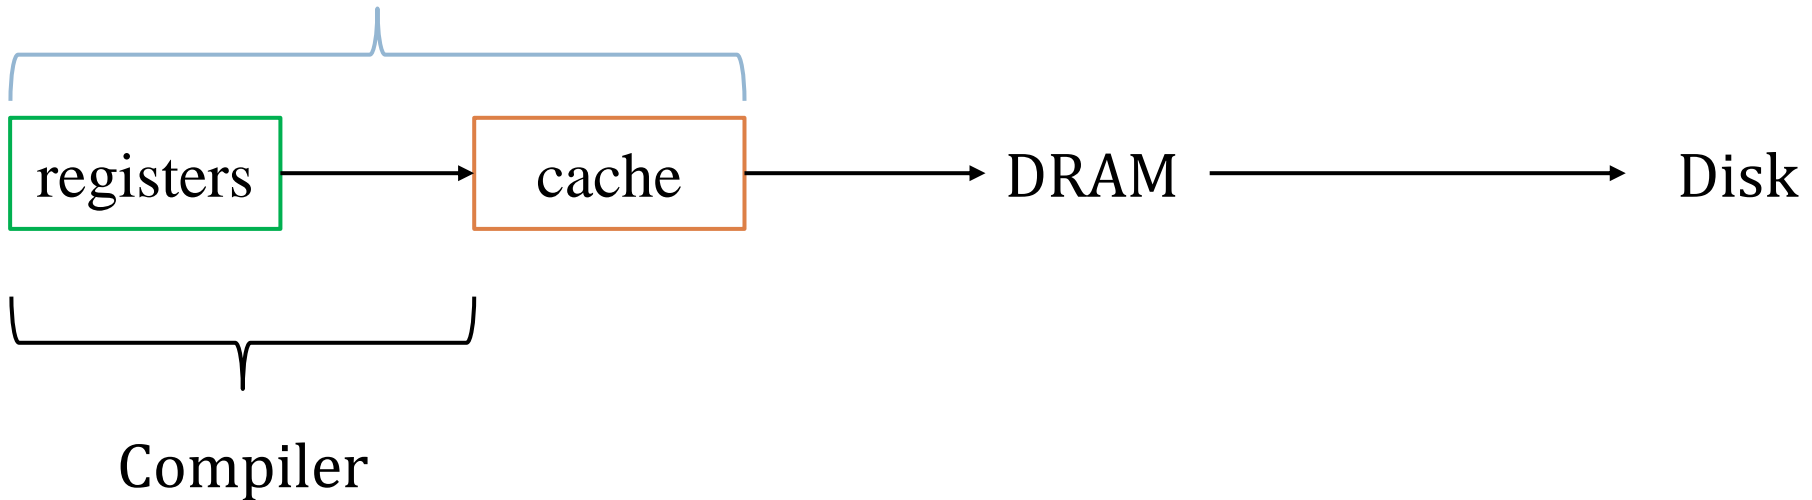

Domains of Data





Domains of Data



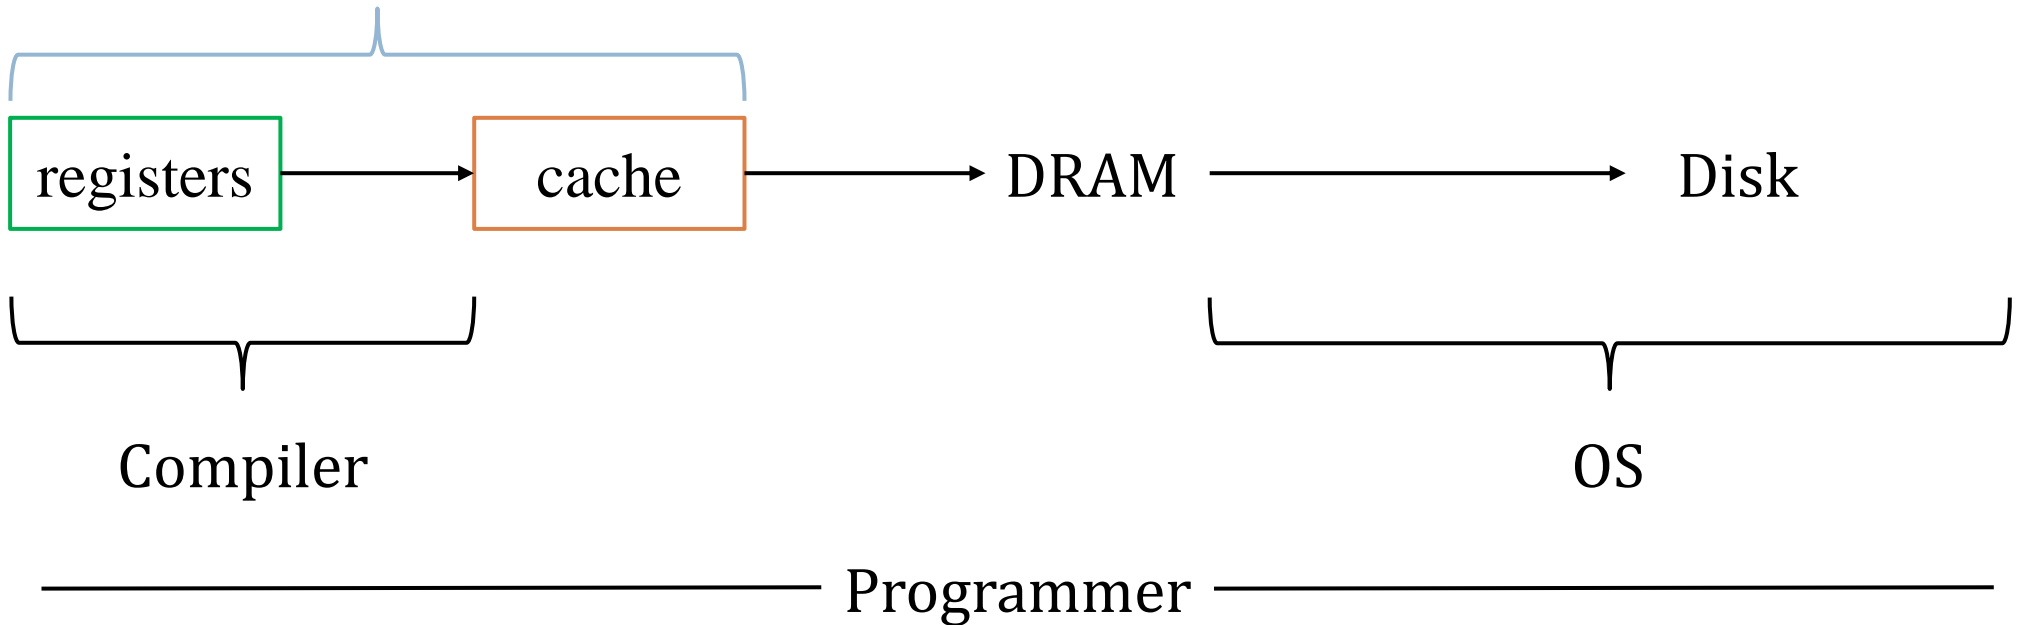

Domains of Data




Domains of Data




Domains of Data




Domains of Data




Domains of Data




Domains of Data



Domains of Data



Programmer



THE ARPANETWORK

DEC 1969

4 NODES

Data Driven

- This end of the spectrum has different challenges:
 - **Vast amounts of data**
 - **Fast access/combination/filtering**
 - **Must be available**
 - Online
 - Securely
 - To simultaneous users

Ponder...

- Suppose you want to save a bunch of Students to a file

- Option 1:

“Jane Doe is a Film major with a GPA of 3.7, and is enrolled in CS2420, and her ID is 12345

John Smith is ...”

Ponder...

- Suppose you want to save a bunch of Students to a file

- Option 1:

“Jane Doe is a Film major with a GPA of 3.7, and is enrolled in CS2420, and her ID is 12345

John Smith is ...”

- How do we search for a student?

Ponder...

- Suppose you want to save a bunch of Students to a file

- Option 1:

“Jane Doe is a Film major with a GPA of 3.7, and is enrolled in CS2420, and her ID is 12345

John Smith is ...”

- How do we search for a student?
 - First we have to know the data's format
 - $O(N)$ scan of entire file

Representing Data

- Option 2: JSON-like (self-describing data)

Major: Film

Class: CS2420

Name: Jane Doe

GPA: 3.7

ID: 12345

Representing Data

- Option 2: JSON-like (self-describing data)

Major: Film

Class: CS2420

Name: Jane Doe

GPA: 3.7

ID: 12345

- How do we find all students enrolled in 2420?
 - **linear scan**

How About XML?

```
<Course>
  <Name>CS2420</Name>
  <Students>
    <Student>
      <Name>Jane Doe</Name>
      <Major>Film</Major>
    </Student>
    <Student>
      <Name>John Smith</Name>
      <Major>CS</Major>
    </Student>
  </Students>
</Course>
```

Still not scalable!

Exercise

- Store a bunch of student records by name, and quickly
 - Add
 - Remove
 - Search / range query
 - Enumerate

Exercise

- Store a bunch of student records by name, and quickly
 - Add
 - Remove
 - Search / range query
 - Enumerate
- Binary search tree

Exercise

- Store a bunch of:
 - students
 - courses
 - professors

Exercise

Professors

Teaching: CS5530, CS4400
Name: Daniel Kopta
ID: 55555

Teaching: CS3500, CS4150
Name: Joe Zachary
ID: 44444

Courses

Name: Database Systems
Num: 5530
Dept. CS

Name: Software Practice
Num: 3500
Dept. CS

Students

Classes: CS5530, Phys2010
Name: Jane Doe
GPA: 3.7
ID: 12345

Classes: CS3500, FILM1010
Name: Jon Smith
GPA: 3.4
ID: 12421

Exercise

Professors

- All courses student *Y* is enrolled in?
- All teachers of student *Z*?
- Order courses by enrollment number?

Students

Classes: CS5530, Phys2010
Name: Jane Doe
GPA: 3.7
ID: 12345

Classes: CS3500, FILM1010
Name: Jon Smith
GPA: 3.4
ID: 12421

Teaching: CS5530, CS4400
Name: Daniel Kopta
ID: 55555

Teaching: CS3500, CS4150
Name: Joe Zachary
ID: 44444

Courses

Name: Database Systems
Num: 5530
Dept. CS

Name: Software Practice
Num: 3500
Dept. CS

Exercise

Professors

Teaching: **CS5530**, CS4400
Name: Daniel Kopta
ID: 55555

Teaching: **CS3500**, CS4150
Name: Joe Zachary
ID: 44444

Courses


Name: Database Systems
Num: **5530**
Dept. **CS**

Name: Software Practice
Num: **3500**
Dept. **CS**

Students

Classes: **CS5530**, Phys2010
Name: Jane Doe
GPA: 3.7
ID: 12345

Classes: **CS3500**, FILM1010
Name: Jon Smith
GPA: 3.4
ID: 12421



Exercise

Professors

Teaching: **CS5530**, CS4400
Name: Daniel Kopta
ID: 55555

Teaching: **CS3500**, CS4150
Name: Joe Zachary
ID: 44444

Courses

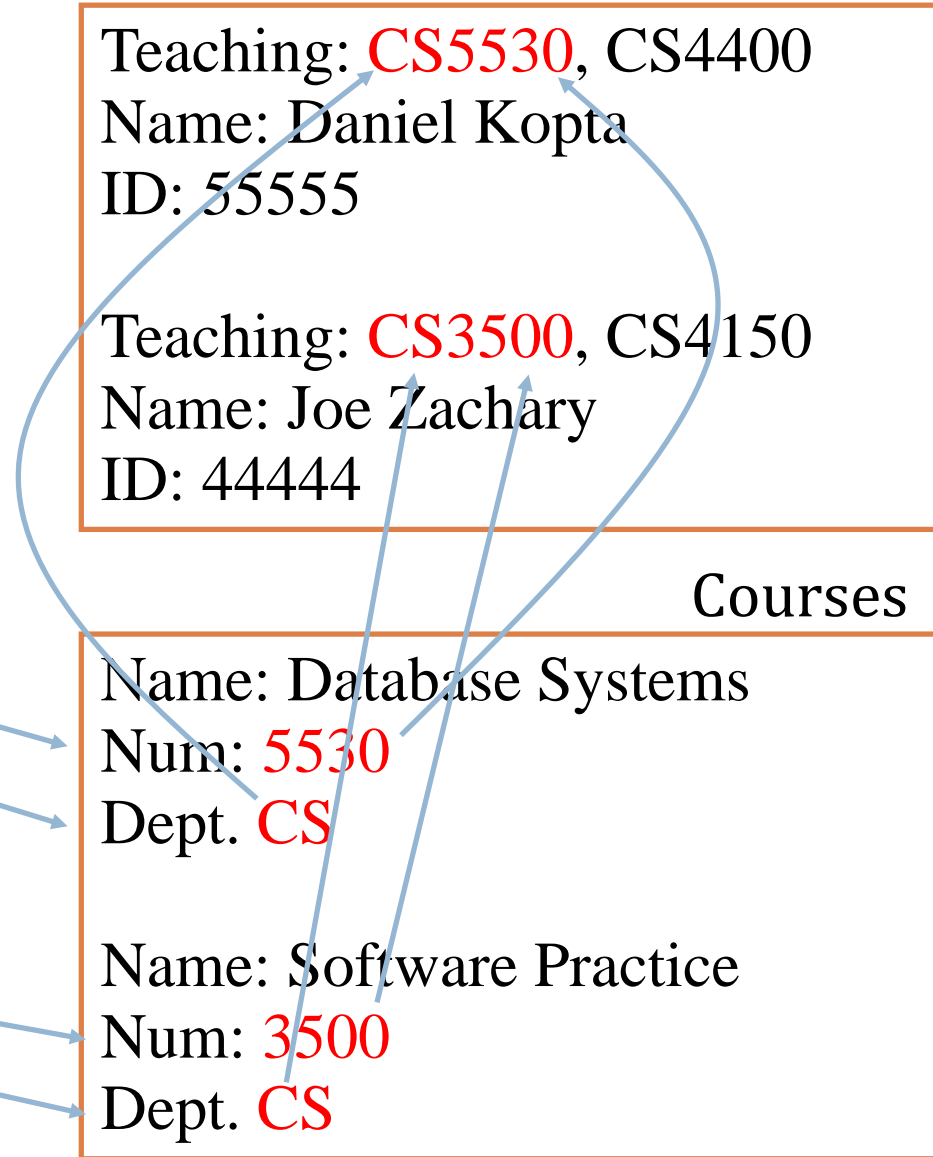
Name: Database Systems
Num: **5530**
Dept. **CS**

Name: Software Practice
Num: **3500**
Dept. **CS**

Students

Classes: **CS5530**, Phys2010
Name: Jane Doe
GPA: 3.7
ID: 12345

Classes: **CS3500**, FILM1010
Name: Jon Smith
GPA: 3.4
ID: 12421



Exercise

- How can we quickly
 - Find all students in course *X*?
 - Find all course(s) student *Y* is enrolled in?
 - Find all teachers of student *Z*?
 - Order students by GPA?
 - Order courses by enrollment number?
 - ...
- Now imagine there are **millions** of each
 - And these operations happen frequently



Solution

1. Structured data
2. Data structures
3. Query language

Solution

- *Structured data*

- Records can not have arbitrary/unpredictable fields/values
- e.g. courses have: **dept**, **num**, and **name**
(string) (int) (string)

Structured Data

- Unstructured

Jane Doe is a Film major with a GPA of 3.7, and her ID is 12345

- Structured

Name (string): “Jane Doe”

Major (string): “Film”

GPA (float): 3.7

ID (uint): 12345

Data Storage

- Save the data itself + data structures
 - Trees, hash tables, etc...

Structured Data



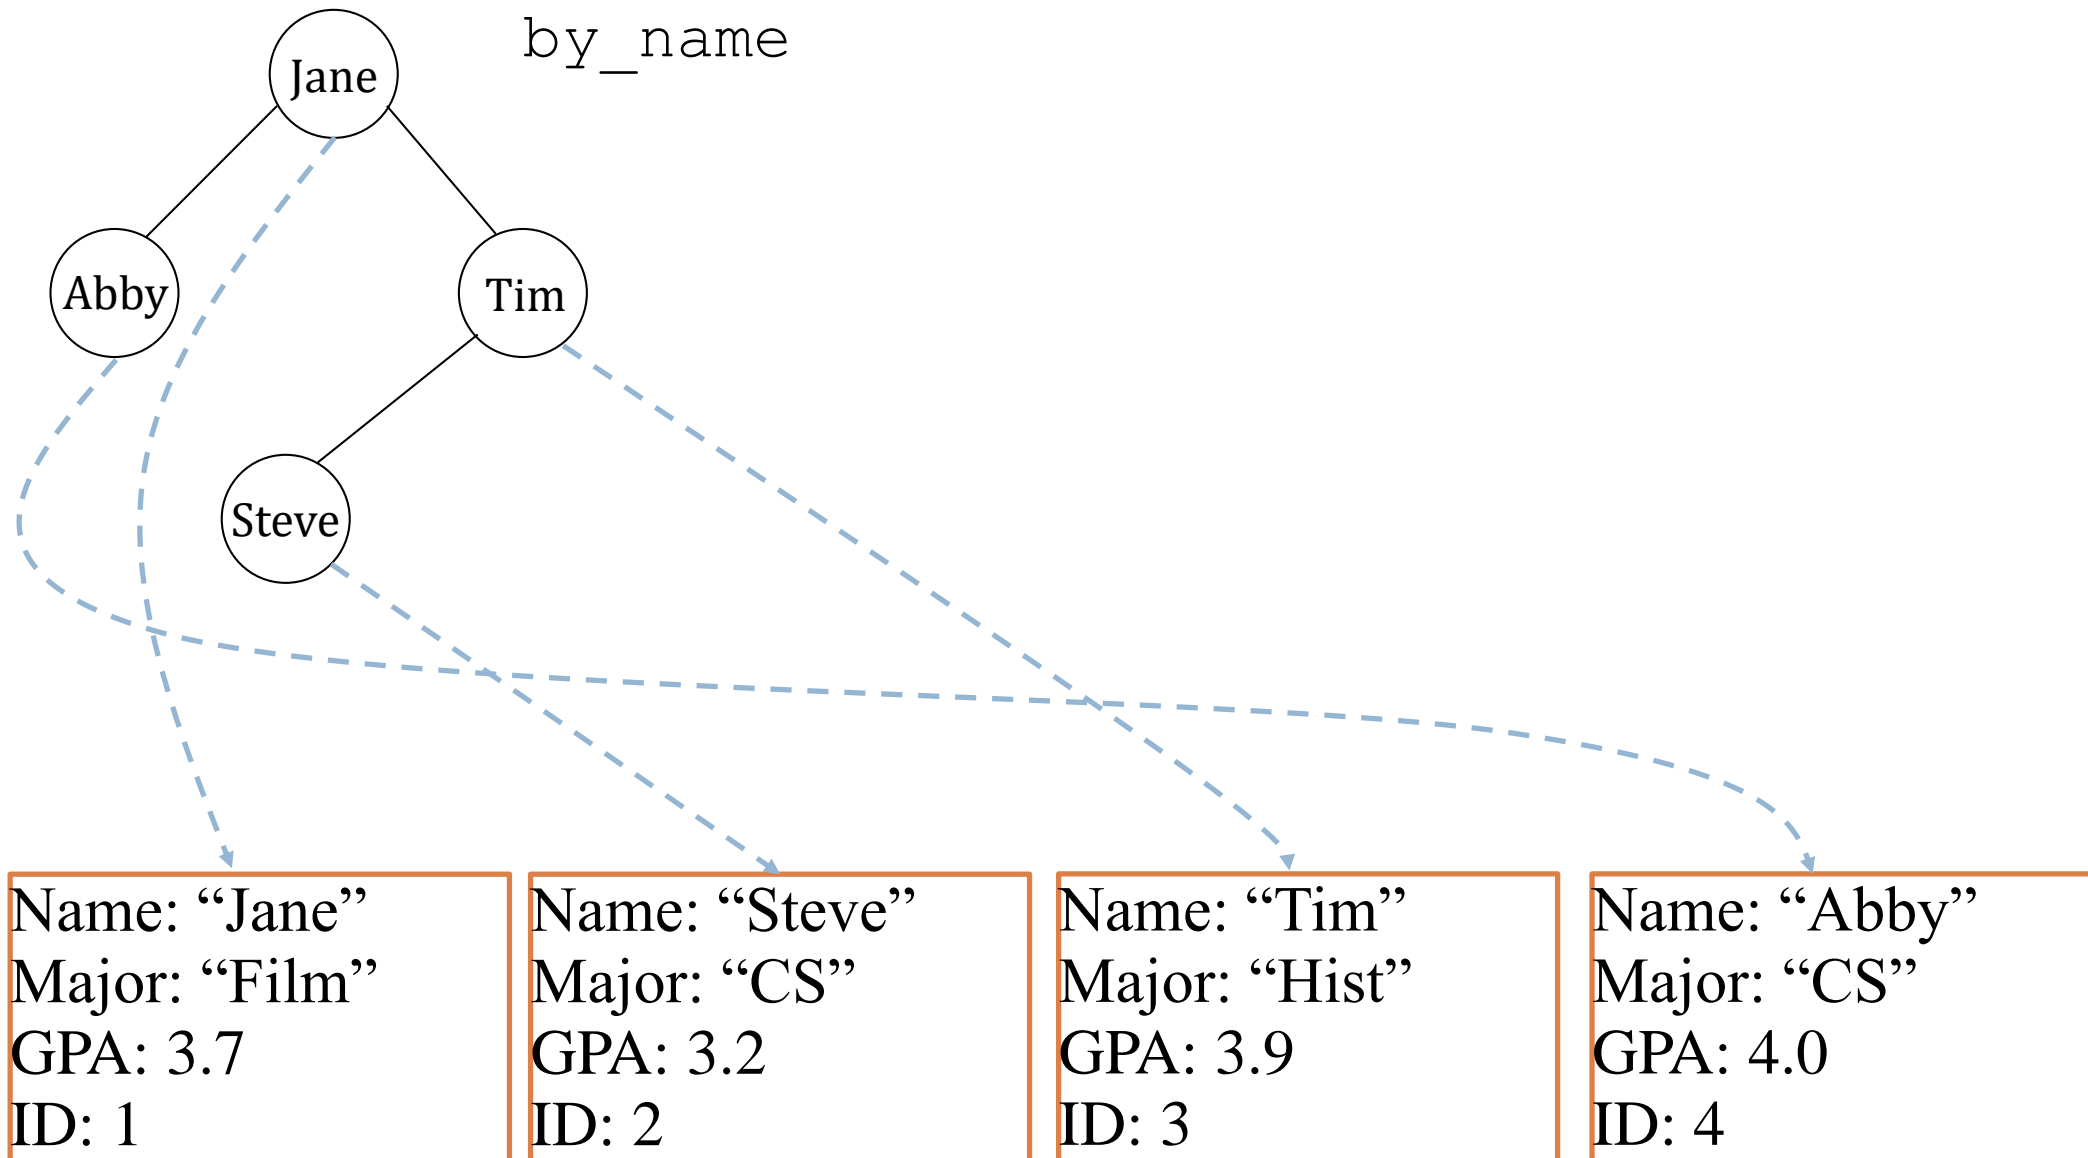
Name: "Jane"
Major: "Film"
GPA: 3.7
ID: 1

Name: "Steve"
Major: "CS"
GPA: 3.2
ID: 2


Name: "Tim"
Major: "Hist"
GPA: 3.9
ID: 3

Name: "Abby"
Major: "CS"
GPA: 4.0
ID: 4

Structured Data + Data Structure



Structured Data + Data Structure



Exercise

- Language for expressing:

- Find all students in course X ?
- Find all course(s) student Y is enrolled in?
- Find all teachers of student Z ?
- Order students by GPA?
- Order courses by enrollment number?
- ...

- C++, Java, C# etc...?

Solution

- Devise language for combining/filtering data

```
SELECT Name FROM Students  
WHERE GPA > 3.5;
```

Solution

1. Structured data
2. Data structures
3. Query language

- ...this is exactly what a *database system* does for you
 - Plus much more!

Why Databases?


- Take advantage of decades of research
 - Availability
 - Reliability
 - Performance
 - Concurrency
 - Interface
- Don't reinvent the wheel

Database System

- Two major components:
 - **Database Management System (DBMS)**
 - Underlying machinery
 - **Query Language**
 - Common interface


Data Storage

Option 1




Data Storage

Option 1




Option 2



Data Storage



Command-line
`mysql> select
* from students`



This Class

- Principles of structured information storage
- Application of databases
- Software interfacing with databases
- Understand DBMS enough to use it effectively

Syllabus

- The syllabus online has been updated – please download a new copy!

Assignments

- A mixture of:
 - Programming
 - Written/diagram
 - Database manipulation

Server

- `atr.eng.utah.edu`
- You all have your own MySQL database on this server

Project

- Implement your own Canvas-like system
 - UI is provided for you
 - You implement the back-end
- Multiple phases throughout the semester



Exams

- One midterm, one final exam
- 5 pages of notes front + back

Grading

- Project: 30%
- Assignments: 30%
- Midterm: 20%
- Final: 20%

Tools

- Visual Studio (Windows)
- C#
- MySQL
- Linux



Visual C#



Visual Studio



Windows

- Windows via Bootcamp recommended
 - Give it at least 80GB
- Parallels *should* also work



Visual Studio

- You will need Visual Studio 2019
 - Don't forget to install .NET
- See Canvas instructions



Class Web Page

- All relevant materials and assignments

utah.instructure.com

- Discuss strategies with other students
- Course announcements
- Grades

Getting Help

- Please ask for help!
- Office hours TBD
- TA consulting hours TBD
- Piazza forums

Academic Misconduct

- Taken very seriously by the SoC
- Any cheating whatsoever results in a failing grade
- Cheating policy is online
- Caught 27 students last year

Relational Databases

- Structured data storage
- Data is organized into **relations**
- Related data are stored “next to” each other
 - e.g. in a table

ID	Name	DOB
...
...

Tables

- Database comprised of one or more tables
- One table represents one relation
 - (pieces of directly-related data)

ID	Name	DOB
1	Harry	31 JUL 1980
2	Hermione	19 SEP 1979
3	Ron	01 MAR 1980
4	Malfoy	05 JUN 1980

Multiple Tables

- Non directly-related data are separated

People

ID	Name	DOB
1	Harry	31 JUL 1980
2	Hermione	19 SEP 1979
3	Ron	01 MAR 1980
4	Malfoy	05 JUN 1980

Courses

Course Num	Name
2420	Alg. and DS
3500	SW Practice
3810	Architecture
4400	Systems
5530	Databases

- Each table is a “relation”

Relation (table)

- Each row is a *tuple* – a set of data units

ID	Name	DOB	GPA
1	Harry	31 JUL 1980	3.5
2	Hermione	19 SEP 1979	4.0
3	Ron	01 MAR 1980	4.0
4	Malfoy	05 JUN 1980	3.9

Relation (table)

- Each row is a *tuple* – a set of data units
 - Does every cell need to be unique?

ID	Name	DOB	GPA
1	Harry	31 JUL 1980	3.5
2	Hermione	19 SEP 1979	4.0
3	Ron	01 MAR 1980	4.0
4	Malfoy	05 JUN 1980	3.9


Relation (table)

- Each row is a *tuple* – a set of data units
 - Does every cell need to be unique? **No**

ID	Name	DOB	GPA
1	Harry	31 JUL 1980	3.5
2	Hermione	19 SEP 1979	4.0
3	Ron	01 MAR 1980	4.0
4	Malfoy	05 JUN 1980	3.9

Relation (table)

- Each row is a *tuple* – a set of data units
 - Does every *row* need to be unique?



ID	Name	DOB	GPA
1	Harry	31 JUL 1980	3.5
2	Hermione	19 SEP 1979	4.0
3	Ron	01 MAR 1980	4.0
4	Malfoy	05 JUN 1980	3.9

Relation (table)

- Each row is a *tuple* – a set of data units
 - Does every *row* need to be unique? **Yes**

ID	Name	DOB	GPA
1	Harry	31 JUL 1980	3.5
2	Hermione	19 SEP 1979	4.0
3	Ron	01 MAR 1980	4.0
4	Malfoy	05 JUN 1980	3.9

Library Example

Name	Phone	CardNum	ISBN	Book
Harry	123-1123	123	978-000	Harry Potter
Hermione	555-1234	124	978-001	A Tale of Two Cities
Ron	123-4567	228	978-002	Last of Us
Malfoy	765-4321	623	978-003	Dune

Library Example

- Is Malfoy directly-related to Dune?

Name	Phone	CardNum	ISBN	Book
Harry	123-1123	123	978-000	Harry Potter
Hermione	555-1234	124	978-001	A Tale of Two Cities
Ron	123-4567	228	978-002	Last of Us
Malfoy	765-4321	623	978-003	Dune

Library Example

- Is Malfoy directly-related to Dune?

Name	Phone	CardNum	ISBN	Book
Harry	123-1123	123	978-000	Harry Potter
Hermione	555-1234	124	978-001	A Tale of Two Cities
Ron	123-4567	228	978-002	Last of Us
Malfoy	765-4321	623	978-003	Dune

- No; only indirectly (he has it checked out)

Library Example

- What if one person checks out multiple books?

Name	Phone	CardNum	ISBN	Book
Harry	123-1123	123	978-000	Harry Potter
Hermione	555-1234	124	978-001	A Tale of Two Cities
Ron	123-4567	228	978-002	Last of Us
Malfoy	765-4321	623	978-003	Dune

Library Example

- What if one person checks out multiple books?

Name	Phone	CardNum	ISBN	Book
Harry	123-1123	123	978-000	Harry Potter
Hermione	555-1234	124	978-001	A Tale of Two Cities
Ron	123-4567	228	978-002	Last of Us
Malfoy	765-4321	623	978-003	Dune
Malfoy	765-4321	623	978-004	Hyperion
Malfoy	765-4321	623	978-005	Bunny Meadows

Library Example

- What if one person checks out multiple books?
 - Duplicate data

Name	Phone	CardNum	ISBN	Book
Harry	123-1123	123	978-000	Harry Potter
Hermione	555-1234	124	978-001	A Tale of Two Cities
Ron	123-4567	228	978-002	Last of Us
Malfoy	765-4321	623	978-003	Dune
Malfoy	765-4321	623	978-004	Hyperion
Malfoy	765-4321	623	978-005	Bunny Meadows

Library Example

- What if one person checks out multiple books?
 - Make a list?

Name	Phone	CardNum	ISBN	Book
Harry	123-1123	123	978-000	Harry Potter
Hermione	555-1234	124	978-001	A Tale of Two Cities
Ron	123-4567	228	978-002	Last of Us
Malfoy	765-4321	623	978-003, 978-004, 978-005	Dune, Hyperion, ...

Library Example

- What if one person checks out multiple books?
 - How to find all people who have checked out “Dune”?
 - How to represent a book that isn’t checked out?
 - How much space to allocate for a row?

Name	Phone	CardNum	ISBN	Book
Harry	123-1123	123	978-000	Harry Potter
Hermione	555-1234	124	978-001	A Tale of Two Cities
Ron	123-4567	228	978-002	Last of Us
Malfoy	765-4321	623	978-003, 978-004, 978-005	Dune, Hyperion, ...

Library Example

- What if one person checks out no books?

Name	Phone	CardNum	ISBN	Book
Harry	123-1123	123	978-000	Harry Potter
Malfoy	765-4321	623	???	???

Even Worse

- Multiple phone numbers, multiple checkouts

Name	Phone	CardNum	ISBN	Title
Dan	888-8888	4	003	Dune
Dan	999-9999	4	003	Dune
Dan	888-8888	4	004	Hyperion
Dan	999-9999	4	004	Hyperion

Solution

- First, let's fix the unrelated-data problem

Patrons

Name	Phone	CardNum
Harry	1 23-1 1 23	1 23
Malfoy	765-4321	623

Inventory

ISBN	Book
978-002	Last of Us
978-003	Dune
978-007	Annihilation

Solution


- First, let's fix the unrelated-data problem
- But what about indirect relationships?
 - How do we specify Malfoy checked out Dune?

Patrons

Name	Phone	CardNum
Harry	123-1123	123
Malfoy	765-4321	623

Inventory

ISBN	Book
978-002	Last of Us
978-003	Dune
978-007	Annihilation



Solution

- Add a table that relates the two

Patrons

Name	Phone	CardNum
Harry	1 23-1 1 23	1 23
Malfoy	765-4321	623

Inventory

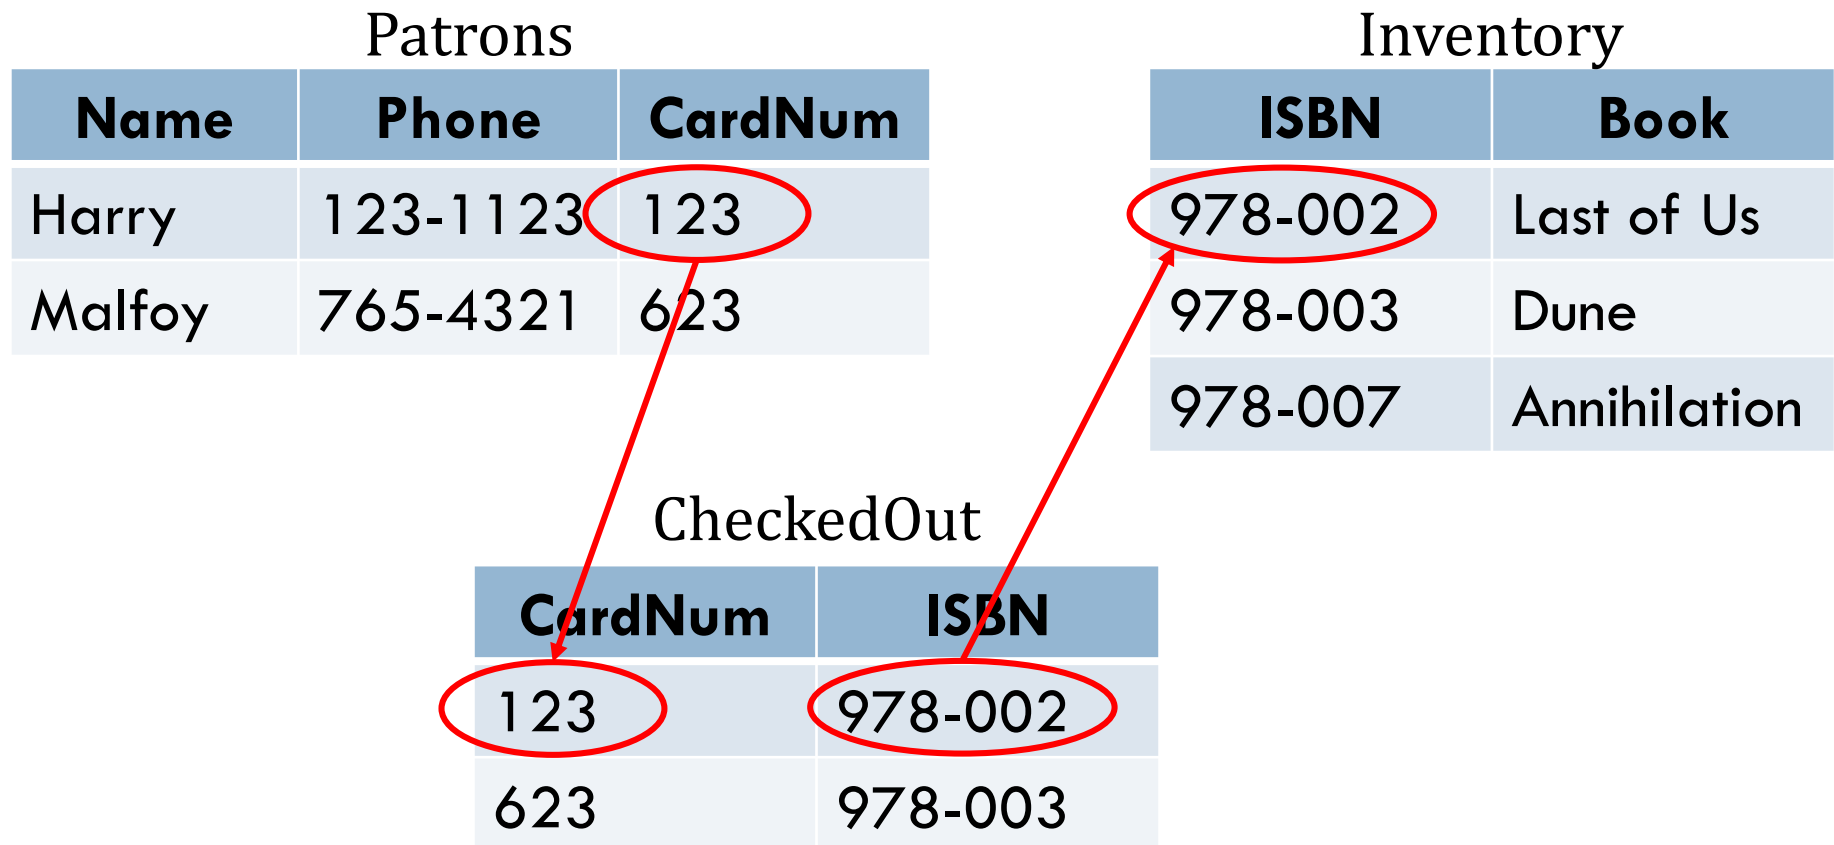
ISBN	Book
978-002	Last of Us
978-003	Dune
978-007	Annihilation

CheckedOut

CardNum	ISBN
1 23	978-002
623	978-003

Solution

- Add a table that relates the two



Solution

- Multiple checkouts
 - Duplicate data minimized

Patrons

Name	Phone	CardNum
Harry	1 23-1 1 23	1 23
Malfoy	765-4321	623

Inventory

ISBN	Book
978-002	Last of Us
978-003	Dune
978-007	Annihilation

CheckedOut

CardNum	ISBN
1 23	978-002
1 23	987-007
623	978-003

Solution

- No checkouts

Patrons

Name	Phone	CardNum
Harry	1 23-1 1 23	1 23
Malfoy	765-4321	623

Inventory

ISBN	Book
978-002	Last of Us
978-003	Dune
978-007	Annihilation

CheckedOut

CardNum	ISBN
623	978-003

Solution

- What about multiple phone numbers?

Patrons

Name	Phone	CardNum
Harry	1 23-1 1 23	1 23
Malfoy	765-4321	623

Inventory

ISBN	Book
978-002	Last of Us
978-003	Dune
978-007	Annihilation

CheckedOut

CardNum	ISBN
623	978-003

Solution

- What about multiple phone numbers?

Patrons

Name	Phone	CardNum
Harry	123-1123	123
Malfoy	765-4321	623

Phones

CardNum	Phone
123	123-1123
123	555-5555
623	765-4321

Inventory

ISBN	Book
978-002	Last of Us
978-003	Dune
978-007	Annihilation

CheckedOut

CardNum	ISBN
623	978-003

Solution

- Pick some unique ID-like field to relate tables (key)
 - **CardNum** and **ISBN**

Patrons

Name	CardNum
Harry	123
Malfoy	623

Phones

CardNum	Phone
123	123-1123
123	555-5555
623	765-4321

Inventory

ISBN	Book
978-002	Last of Us
978-003	Dune
978-007	Annihilation

CheckedOut

CardNum	ISBN
623	978-003

Solution

- Pick some unique ID-like field to relate tables (key)
 - CardNum and ISBN

Patrons

Name	CardNum
Harry	123
Malfoy	623

Inventory


ISBN	Book
978-002	Last of Us
978-003	Dune
978-007	Annihilation

Phones

CardNum	Phone
123	123-1123
123	555-5555
623	765-4321

CheckedOut

CardNum	ISBN
623	978-003



Solution

- What if we have multiple copies of the same book?

Patrons

Name	CardNum
Harry	123
Malfoy	623

Phones

CardNum	Phone
123	123-1123
123	555-5555
623	765-4321

Inventory

ISBN	Book
978-002	Last of Us
978-003	Dune
978-007	Annihilation

CheckedOut

CardNum	ISBN
623	978-003

Solution

- What if we have multiple copies of the same book?
 - **Make another table!**

Patrons

Name	CardNum
Harry	123
Malfoy	623

Phones

CardNum	Phone
123	123-1123
123	555-5555
623	765-4321

Inventory

ISBN	Book
978-002	Last of Us
978-003	Dune
978-007	Annihilation

CheckedOut

CardNum	ISBN
623	978-003

Library

Patrons

Name	CardNum
Joe	1
Ann	2
Ben	3
Dan	4

Inventory

Serial	ISBN
1001	978-0590353427
1002	978-0590353427
1003	978-0679732242
1004	978-0394823379
1005	978-0394823379
1006	978-0062278791

CheckedOut

CardNum	Serial
1	1001
1	1004
4	1005
4	1006

Phones

CardNum	Phone
1	555-5555
2	666-6666
3	777-7777
4	888-8888
4	999-9999

Titles

ISBN	Title	Author
978-0590353427	Harry Potter	Rowling
978-0679732242	The Sound and the Fury	Faulkner
978-0394823379	The Lorax	Seuss
978-0062278791	Profiles in Courage	Kennedy
978-0441172719	Dune	Herbert

Library

Patrons

Name	CardNum
Joe	1
Ann	2
Ben	3
Dan	4

Inventory

Serial	ISBN
1001	978-0590353427
1002	978-0590353427
1003	978-0679732242
1004	978-0394823379

CheckedOut

CardNum	Serial
1	1001
1	1004
4	1005

Phones

CardNum	Phone
1	555-5555
2	666-6666

How do we find all the books checked out by Joe?

Titles

ISBN	Title	Author
978-0590353427	Harry Potter	Rowling
978-0679732242	The Sound and the Fury	Faulkner
978-0394823379	The Lorax	Seuss
978-0062278791	Profiles in Courage	Kennedy
978-0441172719	Dune	Herbert

Library

Patrons

Name	CardNum
Joe	1
Ann	2
Ben	3
Dan	4

Inventory

Serial	ISBN
1001	978-0590353427
1002	978-0590353427
1003	978-0679732242
1004	978-0394823379
1005	978-0394823379
1006	978-0062278791

CheckedOut

CardNum	Serial
1	1001
1	1004
4	1005
4	1006

Phones

CardNum	Phone
1	555-5555
2	666-6666
3	777-7777
4	888-8888
4	999-9999

Titles

ISBN	Title	Author
978-0590353427	Harry Potter	Rowling
978-0679732242	The Sound and the Fury	Faulkner
978-0394823379	The Lorax	Seuss
978-0062278791	Profiles in Courage	Kennedy
978-0441172719	Dune	Herbert

Library

Patrons

Name	CardNum
Joe	1
Ann	2
Ben	3
Dan	4

Inventory

Serial	ISBN
1001	978-0590353427
1002	978-0590353427
1003	978-0679732242
1004	978-0394823379
1005	978-0394823379
1006	978-0062278791

CheckedOut

CardNum	Serial
1	1001
1	1004
4	1005
4	1006

Phones

CardNum	Phone
1	555-5555
2	666-6666
3	777-7777
4	888-8888
4	999-9999

Titles

ISBN	Title	Author
978-0590353427	Harry Potter	Rowling
978-0679732242	The Sound and the Fury	Faulkner
978-0394823379	The Lorax	Seuss
978-0062278791	Profiles in Courage	Kennedy
978-0441172719	Dune	Herbert

Library

Patrons

Name	CardNum
Joe	1
Ann	2
Ben	3
Dan	4

Inventory

Serial	ISBN
1001	978-0590353427
1002	978-0590353427
1003	978-0679732242
1004	978-0394823379
1005	978-0394823379
1006	978-0062278791

CheckedOut

CardNum	Serial
1	1001
1	1004
4	1005
4	1006

Phones

CardNum	Phone
1	555-5555
2	666-6666
3	777-7777
4	888-8888
4	999-9999

Titles

ISBN	Title	Author
978-0590353427	Harry Potter	Rowling
978-0679732242	The Sound and the Fury	Faulkner
978-0394823379	The Lorax	Seuss
978-0062278791	Profiles in Courage	Kennedy
978-0441172719	Dune	Herbert

Library

Patrons

Name	CardNum
Joe	1
Ann	2
Ben	3
Dan	4

Inventory

Serial	ISBN
1001	978-0590353427
1002	978-0590353427
1003	978-0679732242
1004	978-0394823379
1005	978-0394823379
1006	978-0062278791

CheckedOut

CardNum	Serial
1	1001
1	1004
4	1005
4	1006

Phones

CardNum	Phone
1	555-5555
2	666-6666
3	777-7777
4	888-8888
4	999-9999

Titles

ISBN	Title	Author
978-0590353427	Harry Potter	Rowling
978-0679732242	The Sound and the Fury	Faulkner
978-0394823379	The Lorax	Seuss
978-0062278791	Profiles in Courage	Kennedy
978-0441172719	Dune	Herbert

Library

Patrons

Name	CardNum
Joe	1
Ann	2
Ben	3
Dan	4

Phones

CardNum	Phone
1	555-5555
2	666-6666
3	777-7777
4	888-8888
4	999-9999

CheckedOut

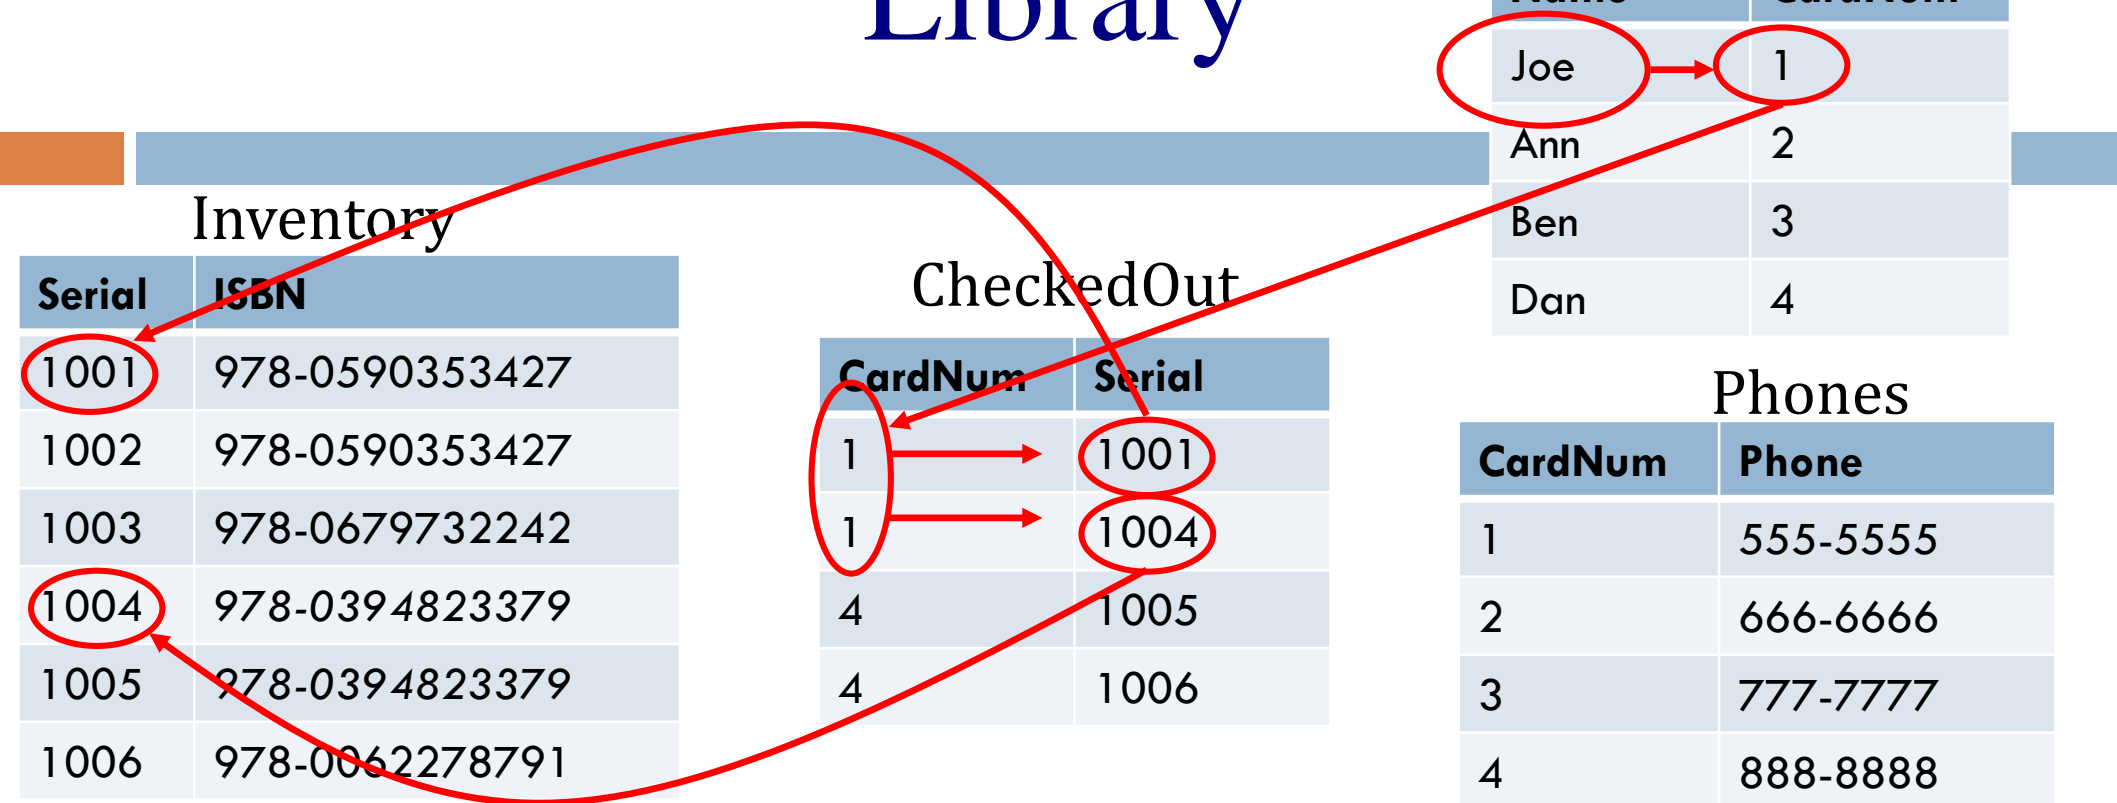
CardNum	Serial
1	1001
1	1004
4	1005
4	1006

Titles

ISBN	Title	Author
978-0590353427	Harry Potter	Rowling
978-0679732242	The Sound and the Fury	Faulkner
978-0394823379	The Lorax	Seuss
978-0062278791	Profiles in Courage	Kennedy
978-0441172719	Dune	Herbert

Inventory

Serial	ISBN
1001	978-0590353427
1002	978-0590353427
1003	978-0679732242
1004	978-0394823379
1005	978-0394823379
1006	978-0062278791



Library

Inventory

Serial	ISBN
1001	978-0590353427
1002	978-0590353427
1003	978-0679732242
1004	978-0394823379
1005	978-0394823379
1006	978-0062278791

CheckedOut

CardNum	Serial
1	1001
1	1004
4	1005
4	1006

Patrons

Name	CardNum
Joe	1
Ann	2
Ben	3
Dan	4

Phones

CardNum	Phone
1	555-5555
2	666-6666
3	777-7777
4	888-8888
4	999-9999

Titles

ISBN	Title	Author
978-0590353427	Harry Potter	Rowling
978-0679732242	The Sound and the Fury	Faulkner
978-0394823379	The Lorax	Seuss
978-0062278791	Profiles in Courage	Kennedy
978-0441172719	Dune	Herbert

Library

Inventory

Serial	ISBN
1001	978-0590353427
1002	978-0590353427
1003	978-0679732242
1004	978-0394823379
1005	978-0394823379
1006	978-0062278791

CheckedOut

CardNum	Serial
1	1001
1	1004
4	1005
4	1006

Patrons

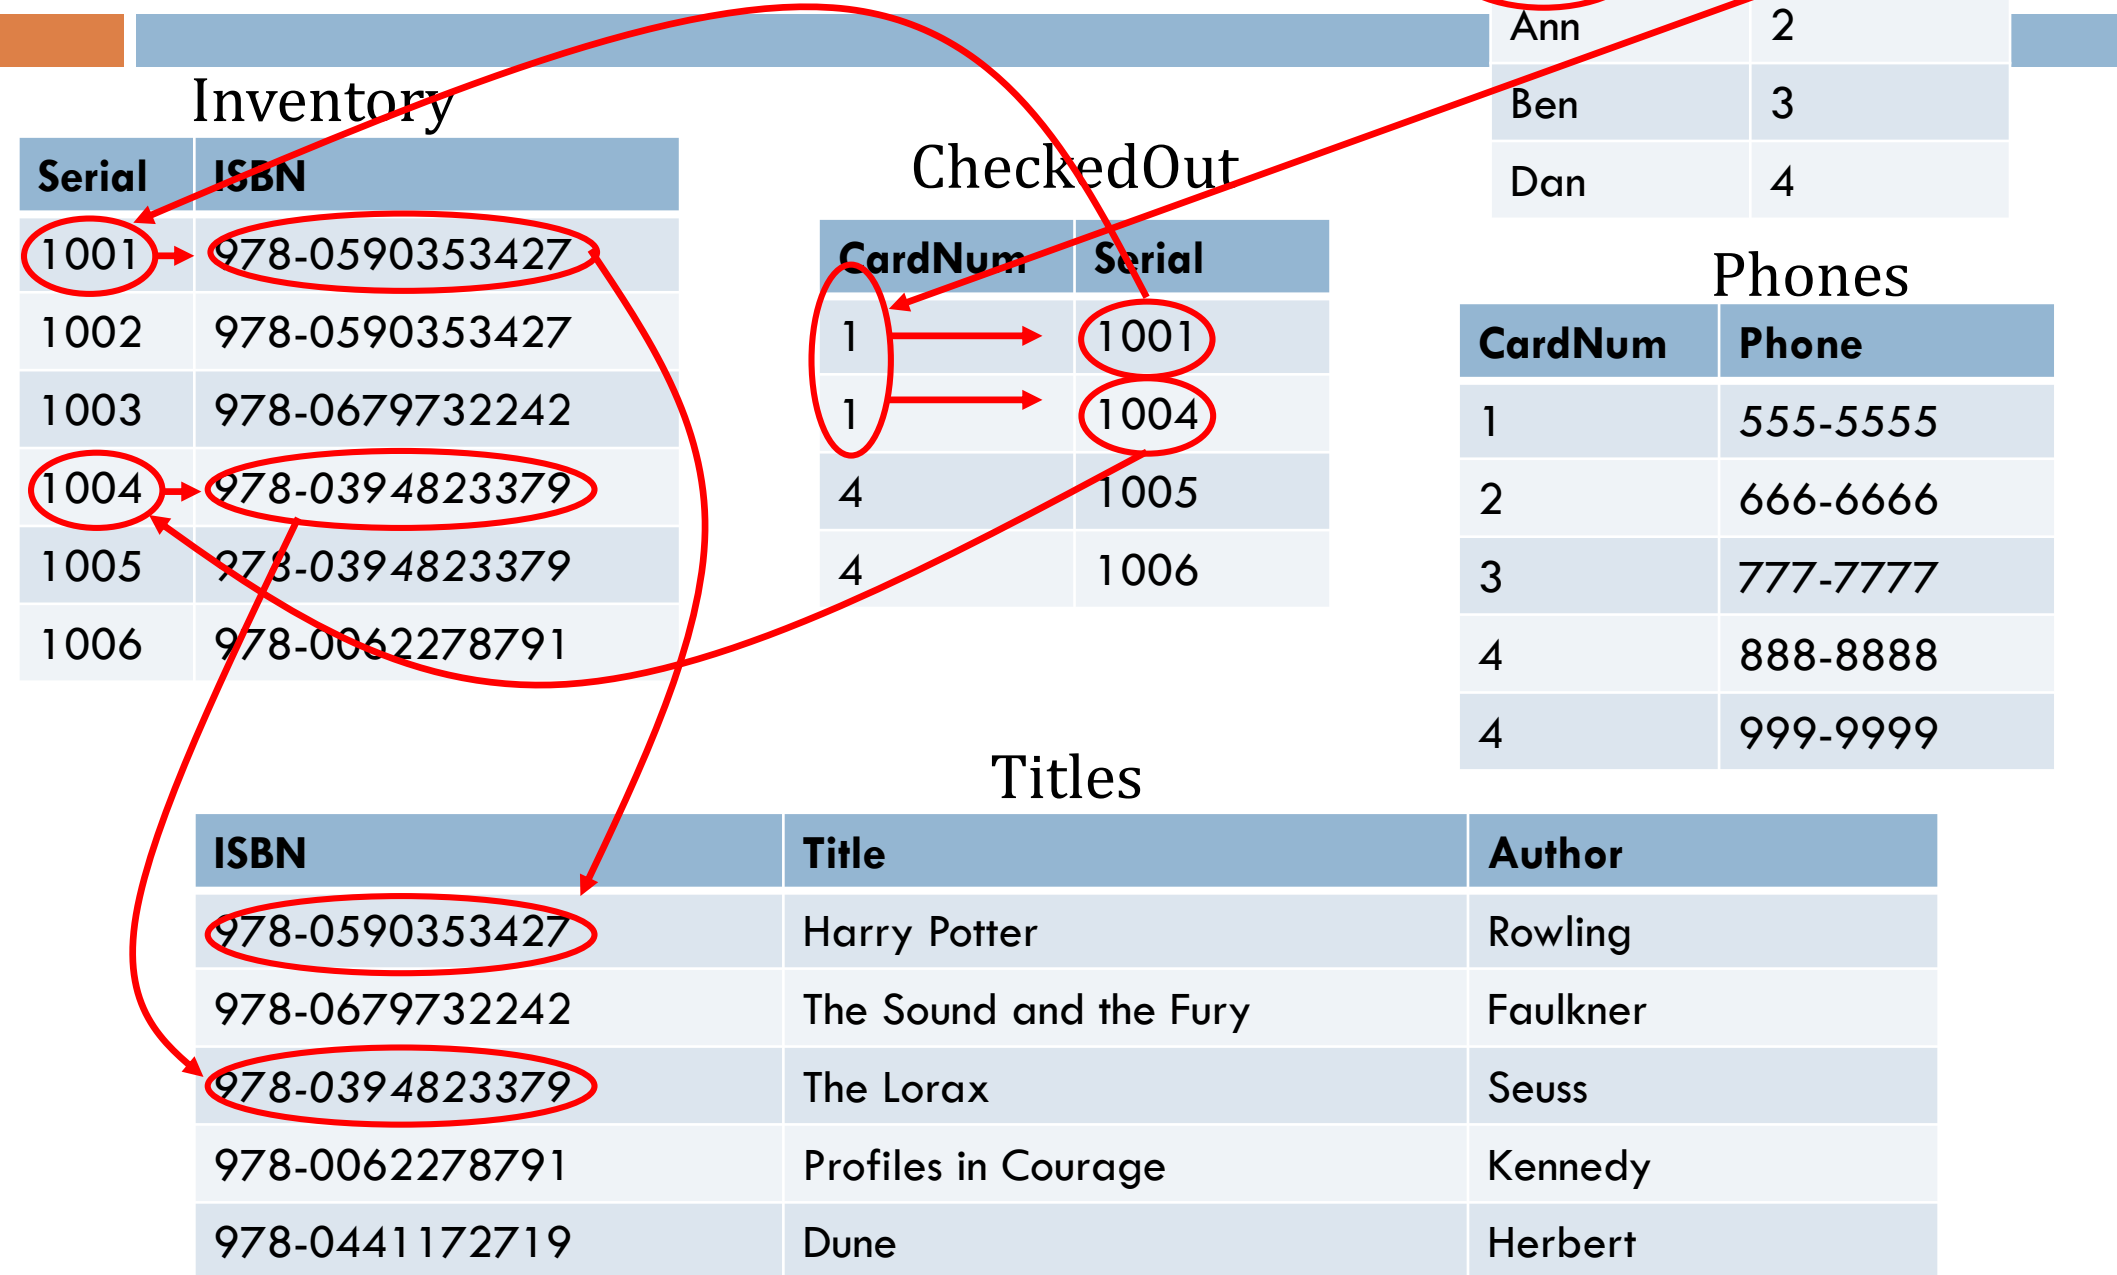
Name	CardNum
Joe	1
Ann	2
Ben	3
Dan	4

Phones

CardNum	Phone
1	555-5555
2	666-6666
3	777-7777
4	888-8888
4	999-9999

Titles

ISBN	Title	Author
978-0590353427	Harry Potter	Rowling
978-0679732242	The Sound and the Fury	Faulkner
978-0394823379	The Lorax	Seuss
978-0062278791	Profiles in Courage	Kennedy
978-0441172719	Dune	Herbert



Library

Inventory

Serial	ISBN
1001	978-0590353427
1002	978-0590353427
1003	978-0679732242
1004	978-0394823379
1005	978-0394823379
1006	978-0062278791

CheckedOut

CardNum	Serial
1	1001
1	1004
4	1005
4	1006

Patrons

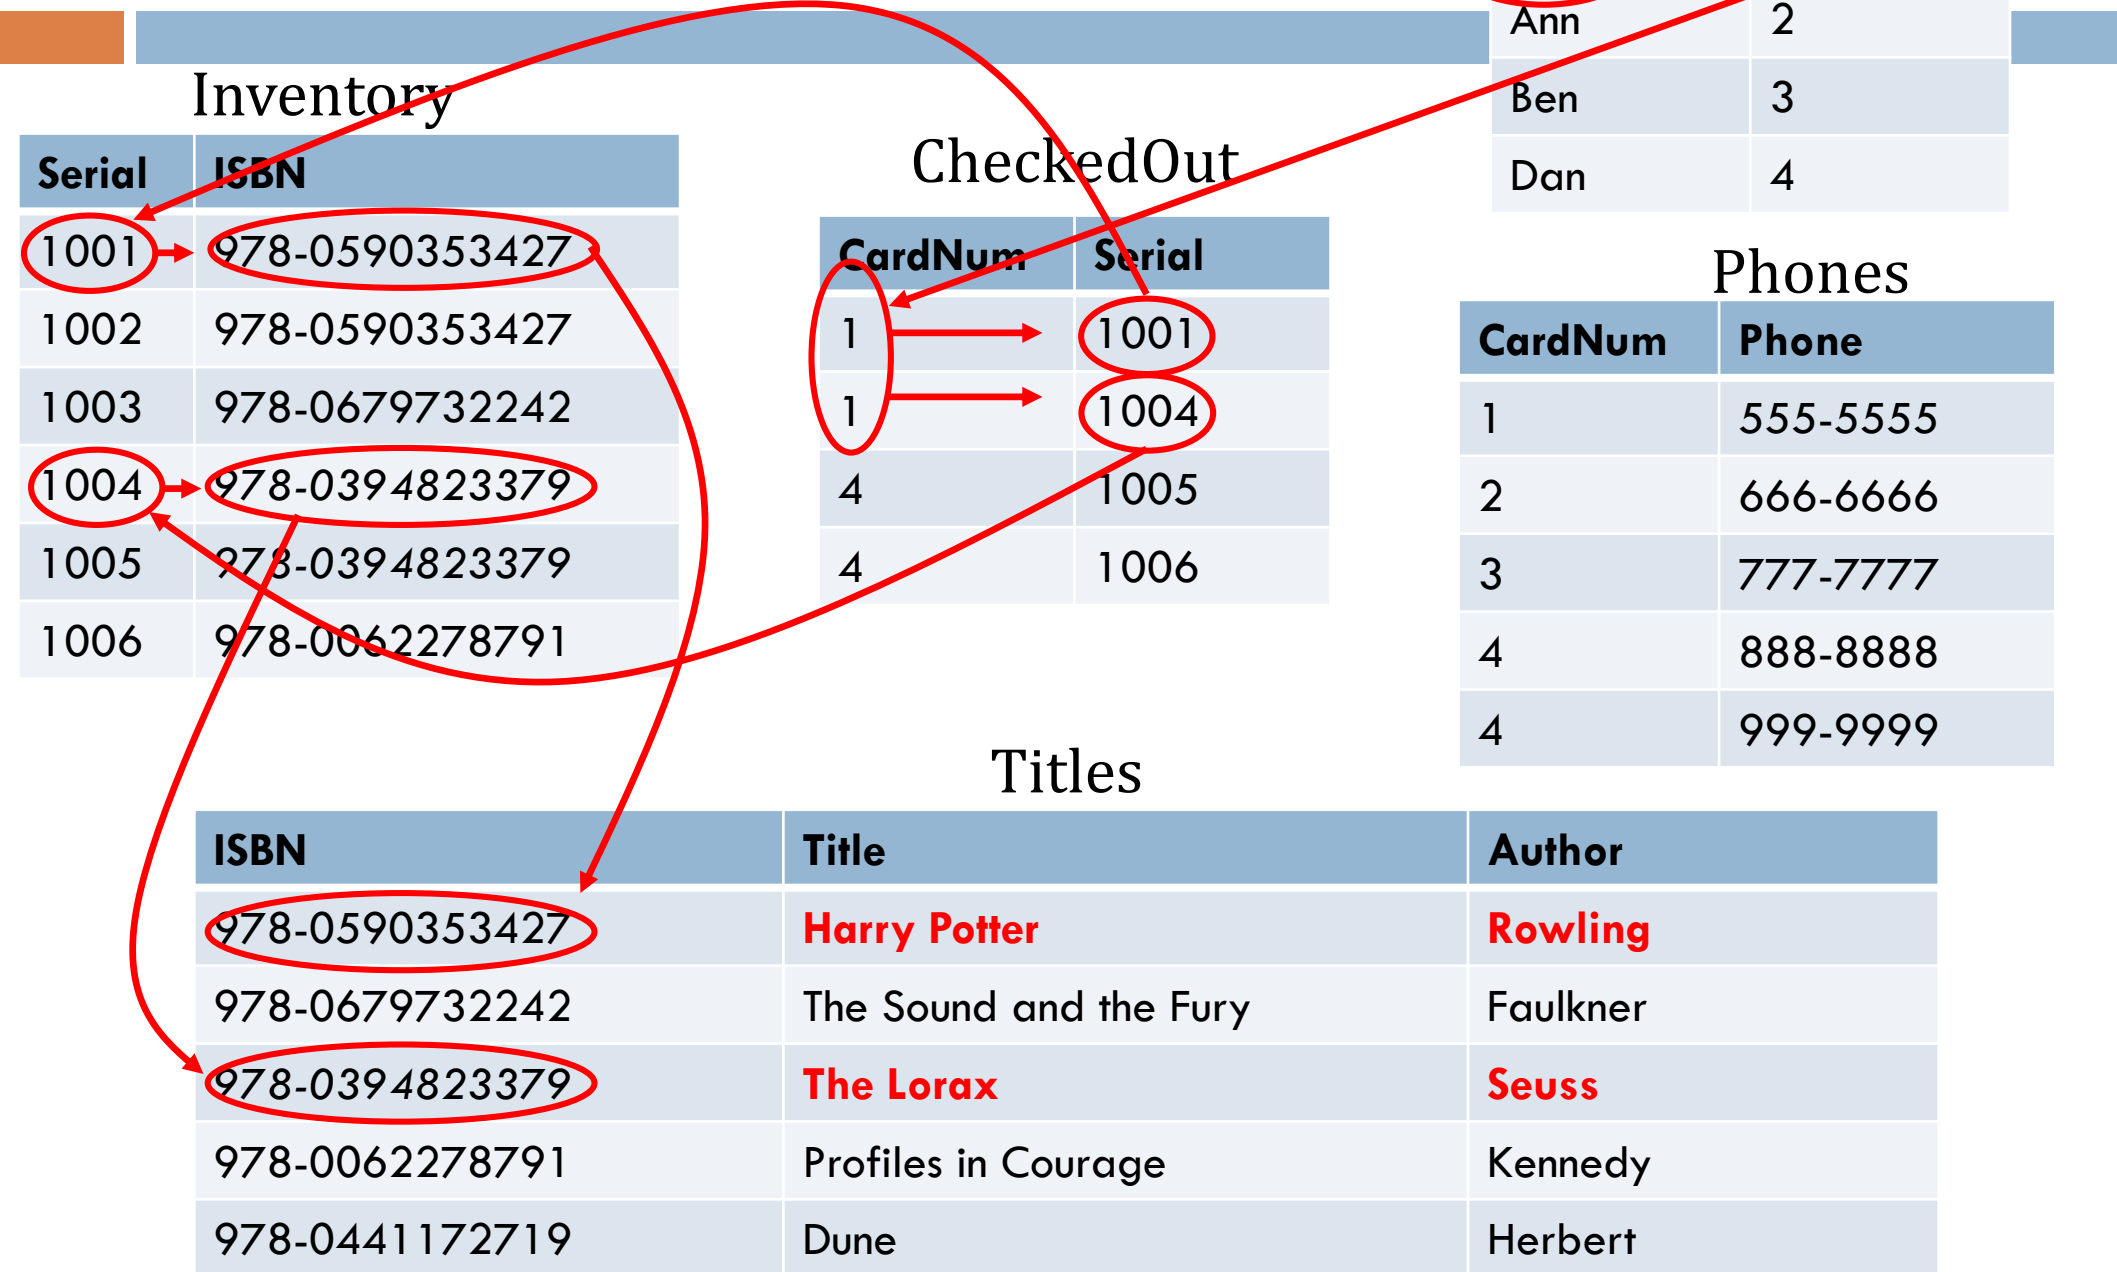
Name	CardNum
Joe	1
Ann	2
Ben	3
Dan	4

Phones

CardNum	Phone
1	555-5555
2	666-6666
3	777-7777
4	888-8888
4	999-9999

Titles

ISBN	Title	Author
978-0590353427	Harry Potter	Rowling
978-0679732242	The Sound and the Fury	Faulkner
978-0394823379	The Lorax	Seuss
978-0062278791	Profiles in Courage	Kennedy
978-0441172719	Dune	Herbert



Don't Worry – It's Fast

Inventory

Serial	ISBN
1001	978-0590353427
1002	978-0590353427
1003	978-0679732242
1004	978-0394823379
1005	978-0394823379
1006	978-0062278791

CheckedOut

CardNum	Serial
1	1001
1	1004
4	1005
4	1006

Patrons

Name	CardNum
Joe	1
Ann	2
Ben	3
Dan	4

Basic Design Goals

- Entries should be atoms (not complex)
 - Don't store lists/arrays

Basic Design Goals

- Entries should be atoms (not complex)
 - Don't store lists/arrays
 - Build compound information by referencing other tables

Basic Design Goals

- Entries should be atoms (not complex)
 - Don't store lists/arrays
 - Build compound information by referencing other tables
 - Enables powerful reasoning about data and relationships, cleaner design

Basic Design Goals

- Entries should be atoms (not complex)
 - Don't store lists/arrays
 - Build compound information by referencing other tables
 - Enables powerful reasoning about data and relationships, cleaner design
 - Enable DBMS to optimize

Basic Design Goals

- Bad news: SQL will let you violate good design rules
- Thus, we design the tables first without even thinking about SQL

Definitions

- **Attribute**: a name and a type (column heading)

Students

Attribute —

sID (uint)	Name (string)	GPA (real)
1	Harry	3.5
2	Hermione	4.0
3	Ron	4.0
4	Malfoy	3.9

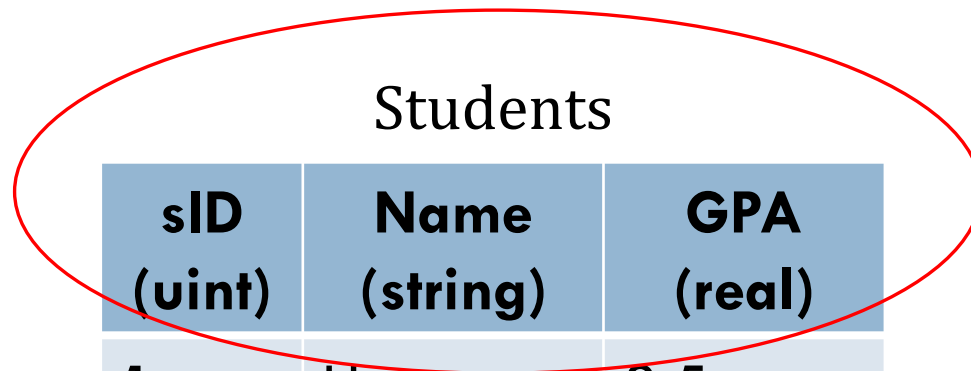
Definitions

- **Schema:** Table name + a set of attributes
 - Specifies the structure/rules of a table

Students

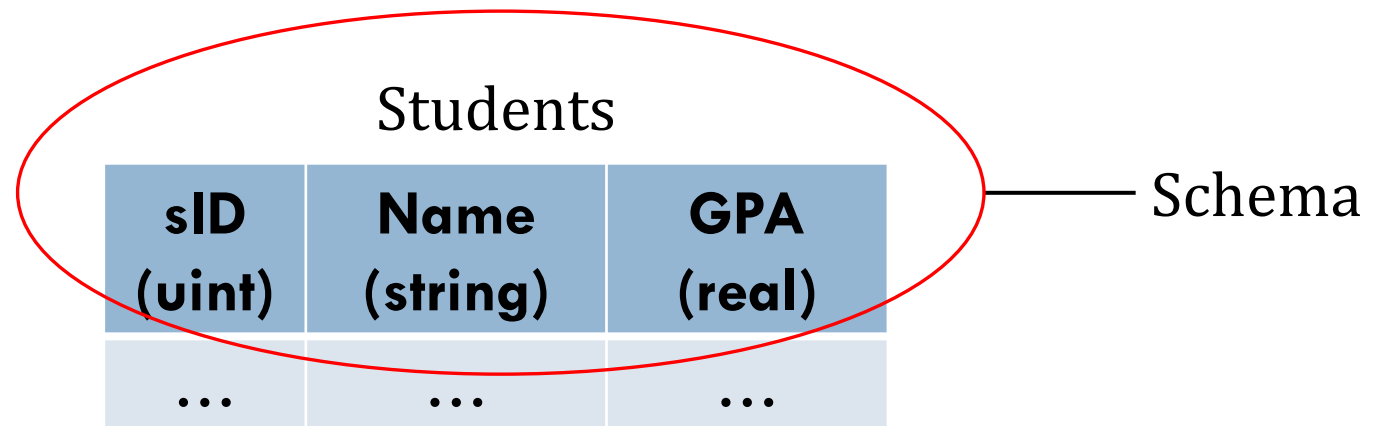
sID (uint)	Name (string)	GPA (real)
1	Harry	3.5
2	Hermione	4.0
3	Ron	4.0
4	Malfoy	3.9

— Schema



Definitions

- **Schema:** Table name + a set of attributes
 - Specifies the structure/rules of a table



- A schema does *not* specify any values

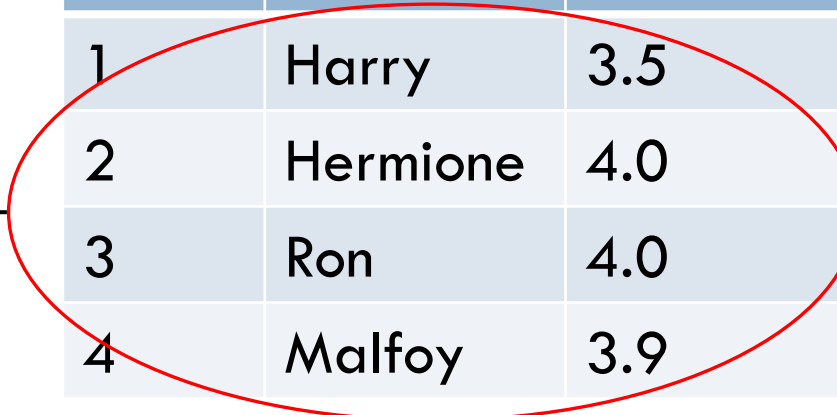
Definitions

- **Instance**: the values in a table
 - A set of *tuples*

Students

sID (uint)	Name (string)	GPA (real)
1	Harry	3.5
2	Hermione	4.0
3	Ron	4.0
4	Malfoy	3.9

Instance —



Definitions

- **Instance**: the values in a table
 - A set of tuples
- **Tuple**: one row

Students

sID (uint)	Name (string)	GPA (real)
1	Harry	3.5
2	Hermione	4.0
3	Ron	4.0
4	Malfoy	3.9

} Tuple

Definitions

- **Relation**: a.k.a “table”
 - Schema + instance

Students

sID (uint)	Name (string)	GPA (real)
1	Harry	3.5
2	Hermione	4.0
3	Ron	4.0
4	Malfoy	3.9

Definitions

- **Relation**: a.k.a “table”
 - A schema + instance

Relation1

sID (uint)	Name (string)	GPA (real)
1	Harry	3.5
2	Hermione	4.0
3	Ron	4.0
4	Malfoy	3.9

Relation2

sID (uint)	Name (string)	GPA (real)
1	Harry	3.5
2	Hermione	4.0

Definitions

- **Instance**: the values in a table
 - A **set** of tuples
- **Tuple**: one row

sID (uint)	Name (string)	GPA (real)
1	Harry	3.5
2	Hermione	4.0
3	Ron	4.0
4	Malfoy	3.9

Definitions

- **Instance**: the values in a table
 - A **set** of tuples – every row is unique!
- **Tuple**: one row

sID (uint)	Name (string)	GPA (real)
1	Harry	3.5
2	Hermione	4.0
3	Ron	4.0
4	Malfoy	3.9

Definitions

- But individual values do not need to be unique
- ...then how do we guarantee each row is unique?

sID (uint)	Name (string)	GPA (real)
1	Harry	3.5
2	Hermione	4.0
3	Ron	4.0
4	Malfoy	3.9

Definitions

- But individual values do not need to be unique
- ...then how do we guarantee each row is unique?
 - **Some attribute (or set of attributes) must be unique**

key

sID (uint)	Name (string)	GPA (real)
1	Harry	3.5
2	Hermione	4.0
3	Ron	4.0
4	Malfoy	3.9

Keys

- Keys uniquely identify each tuple
 - Critical for the DBMS' underlying operations

key

sID (uint)	Name (string)	GPA (real)
1	Harry	3.5
2	Hermione	4.0
3	Ron	4.0
4	Malfoy	3.9

Keys

- Keys uniquely identify each tuple
 - Critical for the DBMS' underlying operations

key		
sID (uint)	Name (string)	GPA (real)
1	Harry	3.5
2	Hermione	4.0
3	Ron	4.0
4	Malfoy	3.9

key		
Make	Model	...
Toyota	Camry	
Toyota	F1	
Subaru	Outback	
Subaru	F1	

Keys

- As a DB designer, you will define keys for each table

Keys

- As a DB designer, you will define keys for each table
- ...but we need formal definitions before SQL