

CS 5530



Database Systems Spring 2020

Relational Model

Keys

Library

Patrons

Name	CardNum
Joe	1
Ann	2
Ben	3
Dan	4

Inventory

Serial	ISBN
1001	978-0590353427
1002	978-0590353427
1003	978-0679732242
1004	978-0394823379
1005	978-0394823379
1006	978-0062278791

CheckedOut

CardNum	Serial
1	1001
1	1004
4	1005
4	1006

Phones

CardNum	Phone
1	555-5555
2	666-6666
3	777-7777
4	888-8888
4	999-9999

Titles

ISBN	Title	Author
978-0590353427	Harry Potter	Rowling
978-0679732242	The Sound and the Fury	Faulkner
978-0394823379	The Lorax	Seuss
978-0062278791	Profiles in Courage	Kennedy
978-0441172719	Dune	Herbert

Library

Patrons

Name	CardNum
Joe	1
Ann	2
Ben	3
Dan	4

Inventory

Serial	ISBN
1001	978-0590353427
1002	978-0590353427
1003	978-0679732242
1004	978-0394823379

CheckedOut

CardNum	Serial
1	1001
1	1004
4	1005

Phones

CardNum	Phone
1	555-5555
2	666-6666

How do we find all the books checked out by Joe?

Titles

ISBN	Title	Author
978-0590353427	Harry Potter	Rowling
978-0679732242	The Sound and the Fury	Faulkner
978-0394823379	The Lorax	Seuss
978-0062278791	Profiles in Courage	Kennedy
978-0441172719	Dune	Herbert

Library

Inventory

Serial	ISBN
1001	978-0590353427
1002	978-0590353427
1003	978-0679732242
1004	978-0394823379
1005	978-0394823379
1006	978-0062278791

CheckedOut

CardNum	Serial
1	1001
1	1004
4	1005
4	1006

Patrons


Name	CardNum
Joe	1
Ann	2
Ben	3
Dan	4


Phones

CardNum	Phone
1	555-5555
2	666-6666
3	777-7777
4	888-8888
4	999-9999

Titles

ISBN	Title	Author
978-0590353427	Harry Potter	Rowling
978-0679732242	The Sound and the Fury	Faulkner
978-0394823379	The Lorax	Seuss
978-0062278791	Profiles in Courage	Kennedy
978-0441172719	Dune	Herbert





Inventory

Serial	ISBN
1001	978-0590353427
1002	978-0590353427
1003	978-0679732242
1004	978-0394823379
1005	978-0394823379
1006	978-0062278791

CheckedOut

CardNum	Serial
1	1001
1	1004
4	1005
4	1006

Patrons

Name	CardNum
Joe	1
Ann	2
Ben	3
Dan	4

Definitions

- **Attribute**: a name and a type (column heading)


Students

Attribute —

sID (uint)	Name (string)	GPA (real)
...

Definitions

- **Schema:** Table name + a set of attributes
 - Specifies the structure/rules of a table



- A schema does *not* specify any values

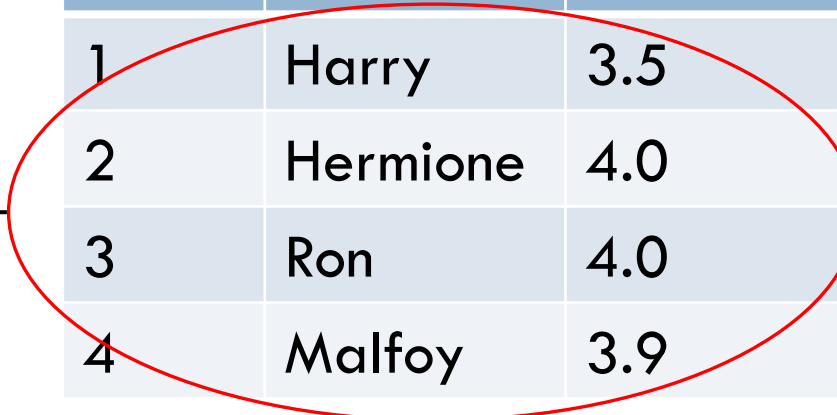
Definitions

- **Instance**: the values in a table
 - A set of *tuples*

Students

sID (uint)	Name (string)	GPA (real)
1	Harry	3.5
2	Hermione	4.0
3	Ron	4.0
4	Malfoy	3.9

Instance —



Definitions

- **Instance**: the values in a table
 - A set of tuples
- **Tuple**: one row

Students

sID (uint)	Name (string)	GPA (real)
1	Harry	3.5
2	Hermione	4.0
3	Ron	4.0
4	Malfoy	3.9

} Tuple

Definitions

- **Relation**: a.k.a “table”
 - Schema + instance

Students

sID (uint)	Name (string)	GPA (real)
1	Harry	3.5
2	Hermione	4.0
3	Ron	4.0
4	Malfoy	3.9

Definitions

- **Relation**: a.k.a “table”
 - A schema + instance

Relation1

sID (uint)	Name (string)	GPA (real)
1	Harry	3.5
2	Hermione	4.0
3	Ron	4.0
4	Malfoy	3.9

Relation2

sID (uint)	Name (string)	GPA (real)
1	Harry	3.5
2	Hermione	4.0

Keys

- Keys uniquely identify each tuple
 - Critical for the DBMS' underlying operations
 - Defines the meaning of the relation

key

sID (uint)	Name (string)	GPA (real)
1	Harry	3.5
2	Hermione	4.0
3	Ron	4.0
4	Malfoy	3.9

key

Make	Model	...
Toyota	F1	
Toyota	Camry	
Subaru	WRX	
Subaru	F1	

Keys

- As a DB designer, you will define keys for each table
- ...but we need formal definitions before SQL

Definitions

- **Superkey:**
- A set of attributes is a *superkey* if no two rows are allowed to have the same values in those columns (no duplicates)

A_1	A_2	A_3
x	4	9
y	4	P
x	3	x

Subsets

A_1

A_2

A_3

A_1, A_2

A_1, A_3

A_2, A_3

A_1, A_2, A_3

{ }

SK?

NO

NO

?

.

?

.

↓

yes

Enrolled

sID	cID	Grade
-----	-----	-------

$\{sID, cID\}$

$\{sID, cID, Grade\}$

Students

sID	Name	GPA
-----	------	-----

$\{sID\}$

$\{sID, \dots, \dots\}$

Definitions

- **Key:**
- A set of fields is a *key* if:
 - it is a *superkey*
 - no proper subset of its fields is a *superkey*

Definitions

- **Key:**
- A set of fields is a *key* if:
 - it is a **superkey**
 - no proper subset of its fields is a **superkey**
- (The empty set is not usually considered a **superkey**)

Keys

- What are the possible key(s)?

$\{A_1, A_2\}$

$\{A_1\}$

$\{A_2\}$

A_1	A_2	A_3
x	4	q
y	4	p
x	3	x

Attribute Set	SK	Key
$\{A_1\}$	No	\longrightarrow
$\{A_2\}$	No	\longrightarrow
$\{A_3\}$	Yes	yes
$\{A_1 A_2\}$	Yes	yes
$\{A_1 A_3\}$	Yes	no
$\{A_2 A_3\}$	Yes	\downarrow
$\{A_1 A_2 A_3\}$	Yes	\downarrow

Keys

- What are the possible key(s)?

A_1	A_2	A_3
x	4	q
y	4	p
x	3	x

Attribute Set	SK	Key
$\{A_1\}$	No	No
$\{A_2\}$	No	No
$\{A_3\}$	Yes	Yes
$\{A_1 A_2\}$	Yes	Yes
$\{A_1 A_3\}$	Yes	No
$\{A_2 A_3\}$	Yes	No
$\{A_1 A_2 A_3\}$	Yes	No

Keys

- DB design works the other way around
 - Start with keys, not data
 - Keys define what data is valid

Keys

- DB designer's job: problem description → schemas

Keys

- What should be the key for the Enrolled schema?

Enrolled		
sID	cID	Grade

- sID = student ID
- cID = course ID
- No instance data given – use human understanding of relation

Keys

- What should be the key for the Enrolled schema?

Enrolled

sID	cID	Grade
------------	------------	--------------

Keys

• What is the key for the Enrolled schema?

- {sID, cID}

Classes

cID

Name when

Students

sID	Name	GPA
1	Harry	3.5
2	Hermione	4.0
3	Ron	4.0
4	Malfoy	3.9

Enrolled

sID	cID	Grade
1	CS3500	B
2	CS3500	A+
4	CS3810	B-
3	CS4400	A-
2	CS6016	A+

Primary/Candidate Key

- **DBA** specifies one key as the **primary key**
- If a schema has more than one key, they are called **candidate keys**

Candidate Keys

- We can specify additional uniqueness constraints (candidate keys)

Enrolled		
sID	cID	Grade

| X A

PK (sID, cID)

UQ (sID, Grade)

Candidate Keys

- We can specify additional uniqueness constraints (candidate keys)

Enrolled		
sID	cID	Grade

PK (sID, cID)

UQ (sID, Grade)

- “A student can not take a course twice, and can not receive the same grade twice”

Exercise

- Provide the primary key (PK) plus any additional candidate keys (UQ) for the `Enrolled` table:

sID	cID	Grade
------------	------------	--------------

1. “A student can not take the same course twice, and no two students can receive the same grade in the same course”
2. “A student can earn multiple grades in a class”
3. “A student can only take one course and receive a single grade for that course, and courses have maximum enrollment of one”

Exercise

- “A student can not take the same course twice, and no two students can receive the same grade in the same course”

$PK(sID, cID)$
 $UQ(cID, Grade)$

sID	cID	Grade
-----	-----	-------

1	X	A
2	X	A

Exercise

- “A student can retake a course for a different grade”

$PK(sID, cID, Grade)$

sID	cID	Grade
------------	------------	--------------

Exercise

- “A student can only take one course and receive a single grade for that course, and courses have a maximum enrollment of one”

$PK(sID)$
 $UK(cID)$ | $PK(sID, cID)$

sID	cID	Grade
-----	-----	-------

Primary vs. Candidate Keys

PK (sID, cID)

UQ (sID, cID)

Or

UQ (sID, Grade)

PK (sID, Grade)

- “A student can not take a course twice, and can not receive the same grade in different courses”

Primary vs. Candidate Keys

PK (sID, cID)

UQ (sID, cID)

Or

UQ (sID, Grade)

PK (sID, Grade)

- “A student can not take a course twice, and can not receive the same grade in different courses”
- PK should be small, ideally integer type(s)

Keys

- Keys relate tuples from different tables

Students

sID	Name	GPA
1	Harry	3.5
2	Hermione	4.0
3	Ron	4.0
4	Malfoy	3.9

Enrolled

sID	cID
1	CS3500
2	CS3500
4	CS3810
3	CS4400
2	CS6016



Quiz

- What happens?

```
Thing* t1 = new Thing();  
Thing* t2 = t1;  
delete(t1);  
cout << t2->x;
```

- a) t2->~~value~~ is printed
- ☒ b) crash
- c) compiler error
- ☒ d) no way to know what happens

Quiz

- What happens?

```
Thing* t1 = new Thing();  
Thing* t2 = t1;  
delete(t1);  
cout << t2->x;
```

- a) t2->value is printed
- b) crash
- c) compiler error
- d) no way to know what happens

Quiz

- What happens?

```
Thing* t1 = new Thing();  
Thing* t2 = t1;  
delete(t1);  
cout << t2->x;
```

- a) t2->value is printed
- b) crash - *what we want to happen (with useful message)*
- c) compiler error
- d) no way to know what happens

Ponder...

- Pointer analogy

Students

sID	Name	GPA
1	Harry	3.5
2	Hermione	4.0
3	Ron	4.0
4	Malfoy	3.9

Enrolled

sID	cID	Grade
4	CS3810	B-
3	CS4400	A-
2	CS6016	A+



Ponder...

- What happens if I delete the student “Ron”?

Students

sID	Name	GPA
1	Harry	3.5
2	Hermione	4.0
3	Ron	4.0
4	Malfoy	3.9

Enrolled

sID	cID	Grade
4	CS3810	B-
3	CS4400	A-
2	CS6016	A+

Ponder...

- What happens if I delete the student “Ron”?

Students

sID	Name	GPA
1	Harry	3.5
2	Hermione	4.0
3	Ron	4.0
4	Malfoy	3.9

Enrolled

sID	cID	Grade
4	CS3810	B-
3	CS4400	A-
2	CS6016	A+

meaningless



Referential Integrity

- **Referential Integrity:**

- References between values should always be meaningful
- This is an *invariant* – should hold true at all times

Referential Integrity

- **Referential Integrity:**

- References between values should always be meaningful
 - This is an *invariant* – should hold true at all times
-
- C++ lets us violate this
 - Java enforces it by taking away control

Referential Integrity

- Option 1:
 - Whenever we remove a record from **Students**, run another command to update all **Enrolled** with that sID

Referential Integrity

- Option 1:

- Whenever we remove a record from **Students**, run another command to update all **Enrolled** with that sID
- Invariant briefly broken
- What if more tables reference **Students**?

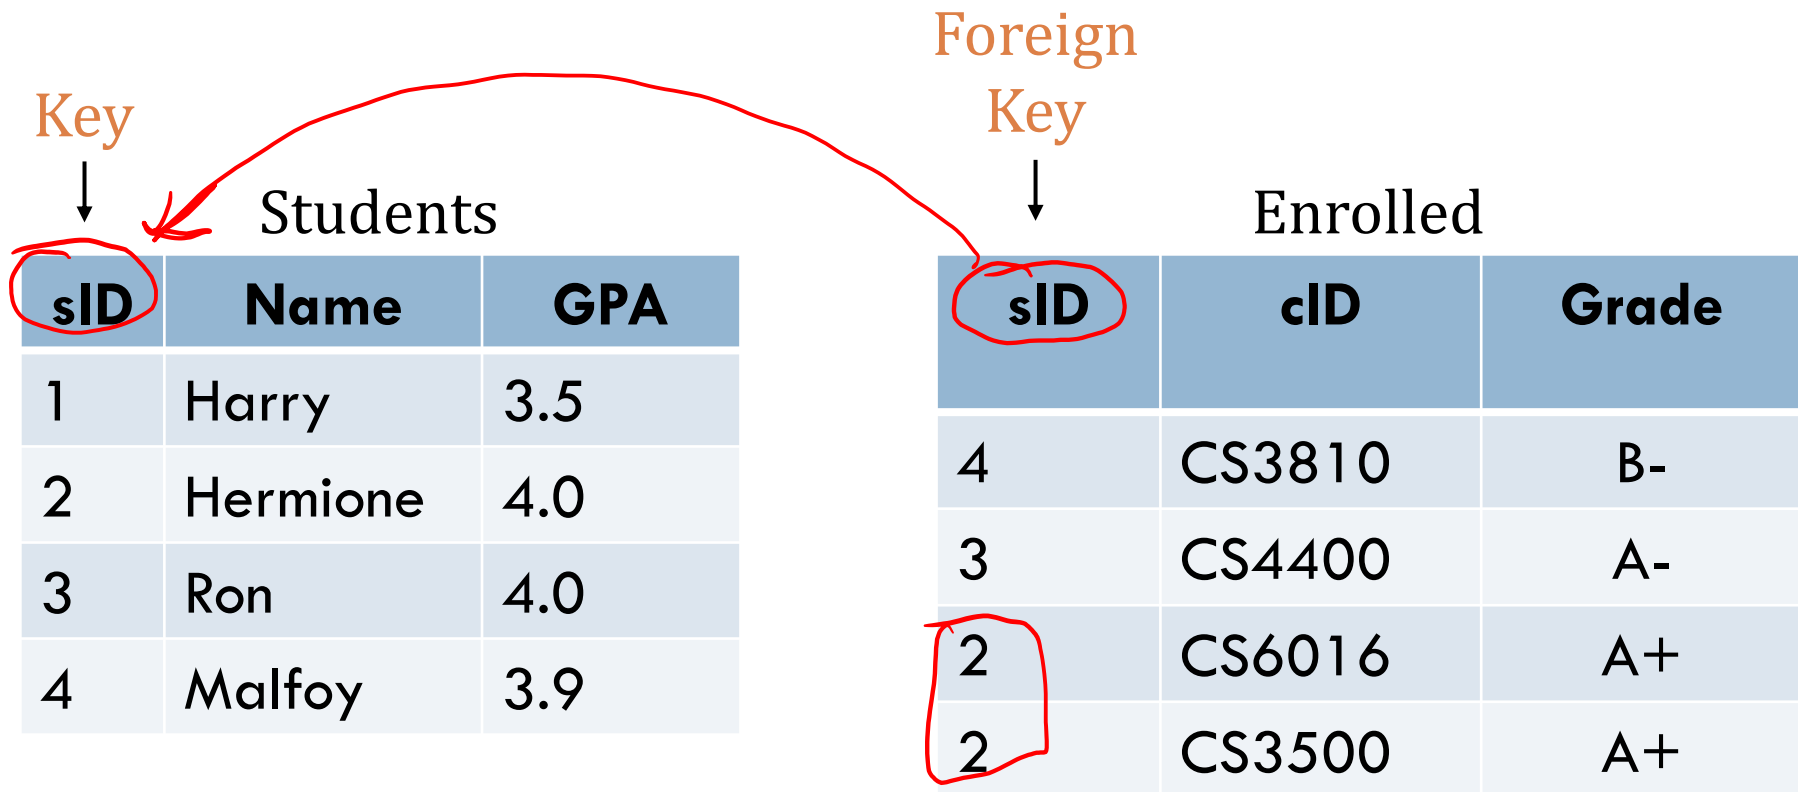
Referential Integrity

- Option 2:
 - Let the DBMS update **Enrolled** automatically
 - SQL provides some support

Foreign Key

- **Foreign Key:**

- Attribute whose values are a key in another table
- Think of it as a “pointer”



Foreign Key

- **Foreign Key:**
 - Not necessarily unique itself

Key



Students

sID	Name	GPA
1	Harry	3.5
2	Hermione	4.0
3	Ron	4.0
4	Malfoy	3.9

Foreign
Key




Enrolled

sID	cID	Grade
4	CS3810	B-
3	CS4400	A-
2	CS6016	A+
2	CS3500	A+

SQL (ish)

```
CREATE TABLE Enrolled(  
  sID int,  
  cID char(6),  
  grade char(2),  
  PRIMARY KEY (sID, cID),  
  FOREIGN KEY (sID) REFERENCES Students(sID) )
```



Students

sID	Name	GPA
1	Harry	3.5
2	Hermione	3.5
3	Ron	4.0
4	Malfoy	3.9

Enrolled

sID	cID	Grade
1	CS3500	A
2	CS3500	B
4	CS3810	B-
3	CS4400	A-

Foreign Key Constraint

- If the *referenced* key is modified, what should we do in the *referencing* table?
- SQL gives a few options

Students

sID	Name	GPA
1	Harry	3.5
2	Hermione	4.0
3	Ron	4.0
4	Malfoy	3.9

Enrolled

sID	cID	Grade
4	CS3810	B-
3	CS4400	A-
2	CS6016	A+

Delete Record

- 1. Delete corresponding record(s)



Patrons

Phones

Name	CardNum
Joe	1
Ann	2
Ben	3
Dan	4

CardNum	Phone
1	555-5555
2	666-6666
3	777-7777
4	888-8888
4	999-9999

Delete Record

- 1. Delete corresponding record(s)

Patrons

Name	CardNum
Joe	1
Ann	2
Ben	3
Dan	4

Phones

CardNum	Phone
1	555-5555
2	666-6666
3	777-7777
4	888-8888
4	999-9999

- If the referencing record has no meaning on its own

Nullify Key

- 2. Nullify foreign key

Students

sID	Name	GPA
1	Harry	3.5
2	Hermione	4.0
3	Ron	4.0
4	Malfoy	3.9

Enrolled

sID	cID	Grade
4	CS3810	B-
NULL	CS4400	A-
2	CS6016	A+

Nullify Key

- 2. Nullify foreign key

Students

sID	Name	GPA
1	Harry	3.5
2	Hermione	4.0
3	Ron	4.0
4	Malfoy	3.9

Enrolled

sID	cID	Grade
4	CS3810	B-
NULL	CS4400	A-
2	CS6016	A+

- If we want to keep the data, but “unlink” it
 - We can still analyze, e.g. average historic GPA

Nullify Key

- 2. Nullify foreign key

Students

sID	Name	GPA
1	Harry	3.5
2	Hermione	4.0
3	Ron	4.0
4	Malfoy	3.9

Enrolled

sID	cID	Grade
4	CS3810	B-
NULL	CS4400	A-
2	CS6016	A+

- But! – This is bad design

- If $\{sID, cID\}$ is primary key for Enrolled

Nullify Key

- 2. Nullify foreign key

Students

sID	Name	GPA
1	Harry	3.5
2	Hermione	4.0
3	Ron	4.0
4	Malfoy	3.9

Enrolled

sID	cID	Grade
4	CS3810	B-
NULL	CS4400	A-
2	CS6016	A+

- But! – This is bad design**

- If {sID, cID} is primary key for Enrolled
- Add a unique ID to Enrolled table

Disallow

- 3. Disallow changes to referenced table
 - Try to delete Joe – SQL reports error

Patrons

Name	CardNum
Joe	1
Ann	2
Ben	3
Dan	4

Steve

5

CheckedOut

CardNum	Serial
1	1001
1	1004
4	1005
4	1006

Disallow

- 3. Disallow changes to referenced table
 - Try to delete Joe – SQL reports error

Patrons

Name	CardNum
Joe	1
Ann	2
Ben	3
Dan	4

CheckedOut

CardNum	Serial
1	1001
1	1004
4	1005
4	1006

- If we need to take some action first
 - Contact Joe, get his books back

Foreign Keys

- What if we try to enroll a non-existent student?
 - SQL will reject it

Students

sID	Name	GPA
1	Harry	3.5
2	Hermione	3.5
3	Ron	4.0
4	Malfoy	3.9



Enrolled

sID	cID	Grade
4	CS3810	B-
3	CS4400	A-
2	CS6016	A
5 (?)	CS4150	E

